

```

from keras.models import Sequential
from keras.layers import Flatten, Dense, Embedding

embedding_dim = 8

model.add(Embedding(vocabulary, embedding_dim, input_length=word_num))

```

$v = 10K$ $d = 8$ $= 20$

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 20, 8)	80000

$\text{word_num} \times \text{embedding_dim}$

```

from keras.models import Sequential
from keras.layers import Flatten, Dense, Embedding

embedding_dim = 8

model = Sequential()
model.add(Embedding(vocabulary, embedding_dim, input_length=word_num))
model.add(Flatten())
model.add(Dense(1, activation='sigmoid'))

model.summary()

```

Layer (type)	Output Shape	Param #
embedding_1 (Embedding)	(None, 20, 8)	80000
flatten_1 (Flatten)	(None, 160)	0
dense_1 (Dense)	(None, 1)	161
Total params: 80,161		
Trainable params: 80,161		
Non-trainable params: 0		

```

from keras import optimizers

epochs = 50

model.compile(optimizer=optimizers.RMSprop(lr=0.0001),
              loss='binary_crossentropy', metrics=['acc'])

```

```

from keras import optimizers

epochs = 50

model.compile(optimizer=optimizers.RMSprop(lr=0.0001),
              loss='binary_crossentropy', metrics=['acc'])
history = model.fit(x_train, y_train, epochs=epochs,
                    batch_size=32, validation_data=(x_valid, y_valid))

```

- The training set is randomly split to a training set and a validation set.
- 80% for training and 20% for validation.
- x_train: $20,000 \times 20$ matrix
- X_valid: $5,000 \times 20$ matrix

Performance on test set

```

loss_and_acc = model.evaluate(x_test, labels_test)
print('loss = ' + str(loss_and_acc[0]))
print('acc = ' + str(loss_and_acc[1]))

```

```

25000/25000 [=====] - 0s 18us/step
loss = 0.5025235502243042
acc = 0.74928

```

- About 75% accuracy on the test set.
- Not bad, because we use only the last 20 words in each movie review. (word_num=20)

Logistic Regression for Sentiment Analysis

seqs[i]:

[27, 28, 29, 30, 31, 22, 32, 33, 34, 35, 1, 36, 29, 37, 38, 39, 40, 2, 41, 42]

10,000×8 parameters

Embedding Layer

X[i]:

word_num × embedding_dim (20×8) matrix

Flatten Layer

x[i]:

160-dim vector

161 parameters

Logistic Regression

f[i]:

Binary Prediction (positive or negative)