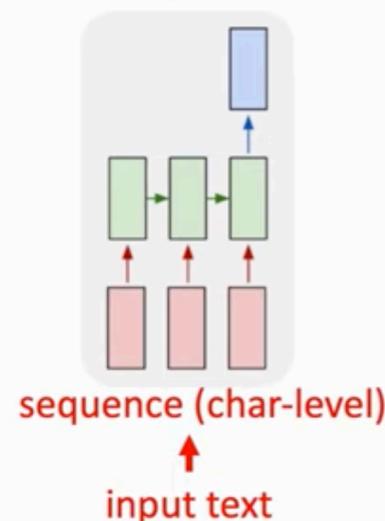


Train an RNN for Text Prediction

predict the next char



How do we train such an RNN?

- Cut text to segments (with overlap).
 - E.g., `seg_len=40` and `stride=3`.

Machine learning is a subset of artificial intelligence in the field of computer science that often uses statistical techniques to give computers the ability to "learn" (i.e., progressively improve performance on a specific task) with data, without being explicitly programmed. The name machine learning was coined in 1959 by Arthur Samuel.

...

...

1. Prepare Training Data

segment:

```
'preface\n\n\nsupposing that truth is a woman--what then? is there  
not ground\nfor suspecting that all philosophers, in so far as they  
have been\nDogmatists, have failed to understand women--that the te  
rrible\nseriousness and clumsy importunity with which they have usu  
ally paid\ntheir addresses to truth, have been unskilled and unseem
```

`segments[0]: preface\n\n\nsupposing that truth is a woman--what then? is there`

`segments[1]: face\n\n\nsupposing that truth is a woman--what then? is there`

`next_chars[0]: i`

`next_chars[1]: a`

1. Prepare Training Data

segment:

```
'preface\n\n\nsupposing that truth is a woman--what then? is there  
not ground\nfor suspecting that all philosophers, in so far as they  
have been\ndogmatists, have failed to understand women--that the te  
rrible\nseriousness and clumsy importunity with which they have usu  
ally paid\ntheir addresses to truth, have been unskilled and unseem
```

`segments[0]: preface\n\n\nsupposing that truth`

`next_chars[0]: i`

`segments[1]: face\n\n\nsupposing that truth is`

`next_chars[1]: a`

`segments[2]: e\n\n\nsupposing that truth is a w`

`next_chars[2]: o`

`segments[3]: \n\nsupposing that truth is a woma`

`next_chars[3]: n`

:

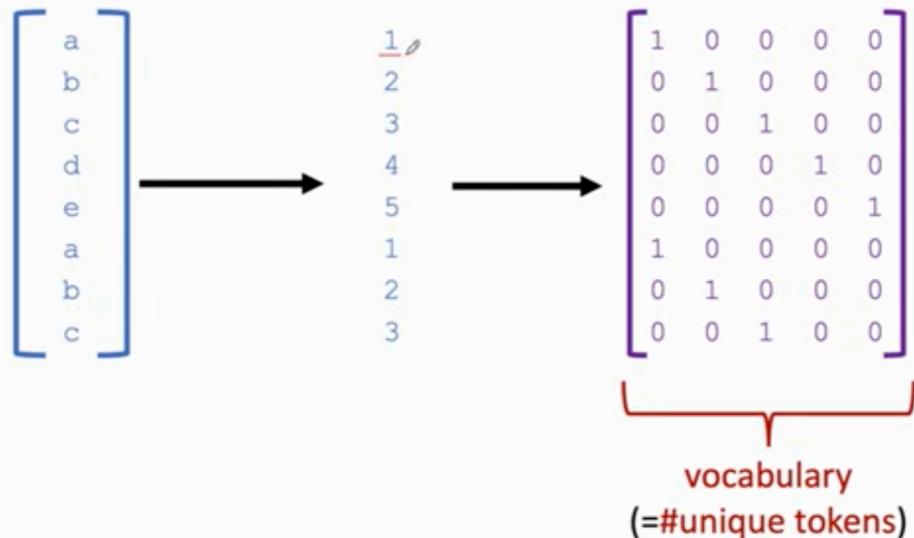
:

2. Character to Vector

Dictionary

Token	Index
a	1
b	2
c	3
d	4
e	5

One-hot encoding (token to vector)



2. Character to Vector

- Vocabulary is $v = 57$ (including letter, blank, punctuation, etc.)
- Segment length is $l = 60$. (A segment has 60 chars.)
- Number of segments is $n = 200,278$. (Number of training samples.)

```
segments[i]: \nsupposing that truth is a woma      next_chars[i]: n  
          ↓  
          X: 60×57 matrix  
          ↓  
          y: 57×1 vector
```

There are $n = 200,278$ such pairs.

3. Build a Neural Network

```
from keras import layers  
  
model = keras.models.Sequential()           l = 60      v = 57  
model.add(layers.LSTM(128, input_shape=(seg_len, vocabulary)))  
model.add(layers.Dense(vocabulary, activation='softmax'))  
model.summary()                          v = 57
```

Layer (type)	Output Shape	Param #
lstm_1 (LSTM)	(None, 128)	95232
dense_1 (Dense)	(None, 57)	7353
Total params: 102,585		
Trainable params: 102,585		
Non-trainable params: 0		

4. Train the Neural Network

```
optimizer = keras.optimizers.RMSprop(lr=0.01)
model.compile(loss='categorical_crossentropy', optimizer=optimizer)
```

```
model.fit(x, y, batch_size=128, epochs=1)
```

```
Epoch 1/1
200278/200278 [=====] - 117s 586us/step -
loss: 1.6746
```

- $x[i, :, :]$ is a segment of 60 chars represented by a 60×57 matrix.
- $y[i, :]$ is the next char represented by a 57-dim vector.

Predict the Next Char

```
pred = model.predict(x_input, verbose=0)[0]
```

Question: Given the **distribution**, what will be the next char?

- Option 1: greedy selection.
 - `next_index = np.argmax(pred)`
 - It is deterministic.
 - Empirically not good.

Predict the Next Char

```
pred = model.predict(x_input, verbose=0)[0]
```

Question: Given the **distribution**, what will be the next char?

- Option 1: greedy selection.
- Option 2: sampling from the multinomial distribution.
 - `next_onehot = np.random.multinomial(1, pred, 1)`
 - `next_index = np.argmax(next_onehot)`
 - Maybe too random.

Predict the Next Char

```
pred = model.predict(x_input, verbose=0)[0]
```

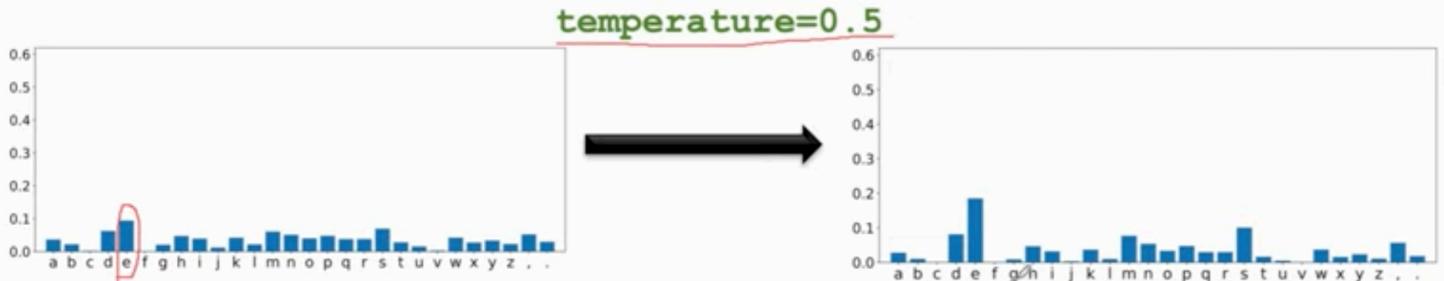
Question: Given the **distribution**, what will be the next char?

- Option 1: greedy selection.
- Option 2: sampling from the multinomial distribution.
- Option 3: adjusting the multinomial distribution.
 - `pred = pred ** (1 / temperature)`
 - `pred = pred / np.sum(pred)`,
 - Sample according to `pred`.
 - Between greedy and multinomial (controlled `temperature`).

Predict the Next Char

Option 3: adjusting the multinomial distribution.

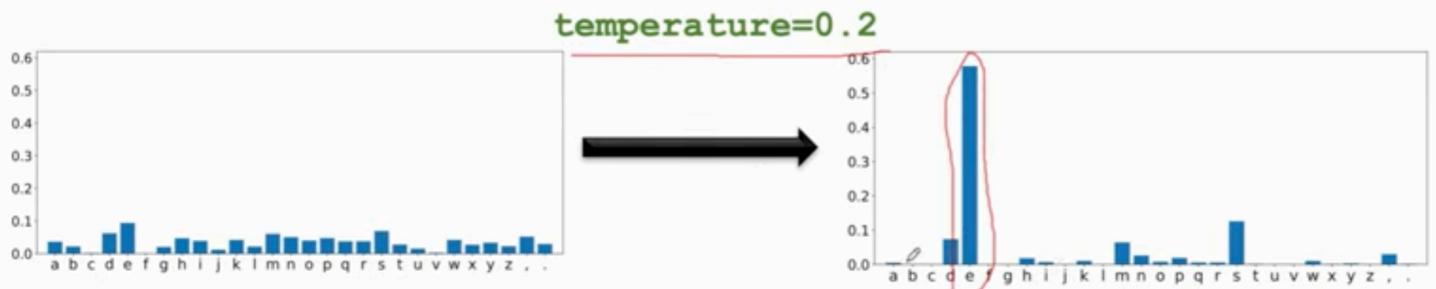
- `pred = pred ** (1 / temperature)`
- `pred = pred / np.sum(pred)`
- Between greedy and multinomial (controlled `temperature`).



Predict the Next Char

Option 3: adjusting the multinomial distribution.

- `pred = pred ** (1 / temperature)`
- `pred = pred / np.sum(pred)`
- Between greedy and multinomial (controlled `temperature`).



Text Generation: An Example

Example: `seg_len=18`

`initial_input (seed): "the cat sat on the"`

`next_char: '_'` (blank)

`input: "he cat sat on the "`

`next_char: 'm'`

`input: "e cat sat on the m"`

Train a Neural Network

1. Partition text to (segment, next_char) pairs.
2. One-hot encode the characters.
 - Character → $v \times 1$ vector.
 - Segment → $l \times v$ matrix.
3. Build and train a neural network.
 - $l \times v$ matrix ==> LSTM ==> Dense ==> $v \times 1$ vector.

Text Generation

1. Propose a seed segment.
2. Repeat the followings:
 - a) Feed the segment (with one-hot) to the neural network.
 - b) The neural network outputs probabilities.
 - c) `next_char` \leftarrow Sample from the probabilities.
 - d) Append `next_char` to the segment.